

## 26.Loan\_predict\_phase3-XGBoost

February 8, 2022

```
[1]: #import packages to do EDA
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

import os
%matplotlib inline
plt.style.use('fivethirtyeight')

import gc

# http://zetcode.com/python/prettytable/
from prettytable import PrettyTable
import matplotlib.font_manager

from sklearn.model_selection import train_test_split
from tqdm import tqdm
import pickle

from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import confusion_matrix, roc_curve, auc
from sklearn.metrics import log_loss
import optuna

import warnings
warnings.filterwarnings("ignore")
```

Read vector from pickle file

```
[2]: with open('./data/train_vector.pkl', 'rb') as f:
      X = pickle.load(f)
```

```
[3]: with open('./data/feature_names.pkl', 'rb') as f:
      feature_names = pickle.load(f)
```

```
[4]: print(X.shape[1],len(feature_names))
```

930 930

```
[5]: with open('./data/yvalues.pkl', 'rb') as f:
      y = pickle.load(f)
```

## 1 testing purpose only

traindata=5000 X = X.toarray()[traindata,:] y=y[traindata] print(X.shape,y.shape)

```
[6]: #split data only train and test.
      #Hypertuning with gridsearch and random
      #hypertuning, will do automatic cv. Hence, split data into Train and Test only.
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33,
      ↪stratify=y,random_state=42)
```

```
[11]: from xgboost import XGBClassifier
```

```
[ ]:
```

## Model Evaluation

### Hyper parameter tuning for ensemble models

```
[7]: #credit:https://www.kaggle.com/saurabhshahane/
      ↪lgbm-hyperparameter-tuning-with-optuna-beginners
      #credit:https://towardsdatascience.com/
      ↪kagglers-guide-to-lightgbm-hyperparameter-tuning-with-optuna-in-2021-ed048d9838b5
      from sklearn.metrics import log_loss
      from sklearn.model_selection import StratifiedKFold
      from optuna import Trial
      from optuna.integration import XGBoostPruningCallback

      from scipy.stats import randint as sp_randint
      from sklearn.metrics import roc_auc_score

      def objective(trial, X_data, y_label):
          param_grid = {
              "verbosity": 0,
              "objective" : trial.suggest_categorical("objective", ['binary:
              ↪logistic']),
              "n_estimators": trial.suggest_int('n_estimators', 50, 10000),
              "learning_rate": trial.suggest_categorical('learning_rate', [0.008,0.
              ↪009,0.01,0.012,0.014,0.016,0.018, 0.02]),
              "alpha": trial.suggest_loguniform('alpha', 1e-3, 10.0),
              'gamma': trial.suggest_loguniform('gamma', 1e-3, 10.0),
              "lambda": trial.suggest_loguniform('lambda', 1e-3, 10.0),
```

```

        "min_child_weight": trial.suggest_int('min_child_weight', 1, 200),
        'subsample' : trial.suggest_loguniform('subsample', 0.5, 1),
        'colsample_bytree': trial.suggest_loguniform('colsample_bytree', 0.5,
→1),
        'colsample_bylevel': trial.suggest_loguniform('colsample_bylevel', 0.5,
→1),
        "max_depth": trial.suggest_categorical('max_depth',
→[5,7,9,11,13,15,17,20]),
    }

    cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    cv_scores = np.empty(5)

    for idx, (train_idx, valid_idx) in enumerate(cv.split(X_data, y_label)):
        x_train, x_valid = X_data[train_idx], X_data[valid_idx]
        y_train, y_valid = y_label[train_idx], y_label[valid_idx]

        #Define weight ratio
        weight_ratio = float(len(y_train[y_train == 0]))/
→float(len(y_train[y_train ==1]))
        w_array = np.array([1]*y_train.shape[0])#positive class
        w_array[y_train==1] = weight_ratio
        w_array[y_train==0] = 1- weight_ratio #negative class

        #XGBoostClassifier with random_state=0
        #scale_pos_weight - A value greater than 0 should be used
        #in case of high class imbalance as it helps in faster convergence

        model = XGBClassifier(tree_method = "exact",predictor = "cpu_predictor",
→
→scale_pos_weight=(weight_ratio*100),random_state=0)
        model.set_params(**param_grid)

        model.fit(x_train,y_train, eval_set=[(x_valid, y_valid)], verbose=0,
→
→early_stopping_rounds=50,callbacks=[XGBoostPruningCallback(trial,
→'validation_0-logloss')]
        )

        preds = model.predict_proba(x_valid)
        cv_scores[idx] = roc_auc_score(y_valid, preds[:,1])

    return np.mean(cv_scores)

```

```

[8]: from warnings import simplefilter
    simplefilter("ignore", category=RuntimeWarning)

```

```

optuna.logging.set_verbosity(optuna.logging.WARNING)
study = optuna.create_study(direction="minimize", study_name="XGBoost_
    ↳Classifier" )
func = lambda trial: objective(trial, X_train, y_train)
study.optimize(func, n_trials=100)

```

```

[9]: print(f"\tBest value (rmse): {study.best_value:.5f}")
      print(f"\tBest params:")

      for key, value in study.best_params.items():
          print(f"\t\t{key}: {value}")

```

```

Best value (rmse): 0.66379
Best params:
    objective: binary:logistic
    n_estimators: 5509
    learning_rate: 0.008
    alpha: 0.0014764404145375422
    gamma: 0.006709171934295603
    lambda: 0.004841559984476445
    min_child_weight: 18
    subsample: 0.801485530162779
    colsample_bytree: 0.8775747304615185
    colsample_bylevel: 0.5456219632677068
    max_depth: 20

```

```

[42]: def batch_predict(clf, data):
        # roc_auc_score(y_true, y_score) the 2nd parameter should be probability_
        ↳estimates of the positive class
        # not the predicted outputs
        y_data_pred = []
        tr_loop = data.shape[0] - data.shape[0]%1000
        # consider you X_tr shape is 49041, then your tr_loop will be 49041 -
        ↳49041%1000 = 49000
        # in this for loop we will iterate until the last 1000 multiplier
        for i in range(0, tr_loop, 1000):
            y_data_pred.extend(clf.predict_proba(data[i:i+1000]))[:,1])
        # we will be predicting for the last data points
        if data.shape[0]%1000 !=0:
            y_data_pred.extend(clf.predict_proba(data[tr_loop:]))[:,1])
        return y_data_pred

```

```

[43]: #https://stackoverflow.com/questions/61748441/
        ↳how-to-fix-the-values-displayed-in-a-confusion-matrix-in-exponential-form-to-normal
        def plot_confusionmatrix(y_tr,y_trpred,y_te,y_tepred):
            from sklearn.metrics import confusion_matrix

```

```

tn, fp, fn, tp = confusion_matrix(y_tr, np.round(y_trpred)).ravel()
print('Traing data tn-> {}, fp-> {}, fn-> {}, tp-> {}'.format(tn, fp, fn,
→tp), end=" ")
#confusion matrix on training data
plt.figure(figsize=(10, 10))
ax_tr = plt.subplot(221)
cm_tr = confusion_matrix(y_tr, np.round(y_trpred))
plt.title("Training data - Confusion Matrix")
sns.heatmap(cm_tr, ax=ax_tr, fmt='d', cmap='YlGnBu', annot=True)
# labels, title and ticks
ax_tr.set_xlabel('Predicted labels');
ax_tr.set_ylabel('True labels');
ax_tr.set_ylim(2.0, 0)
ax_tr.xaxis.set_ticklabels(['No', 'Yes']);
ax_tr.yaxis.set_ticklabels(['No', 'Yes']);

#Confusion matrix on test data
tn, fp, fn, tp = confusion_matrix(y_te, np.round(y_tepred)).ravel()
print('Traing data tn-> {}, fp-> {}, fn-> {}, tp-> {}'.format(tn, fp, fn,
→tp), end=" ")

ax_te = plt.subplot(222)
cm_te = confusion_matrix(y_te, np.round(y_tepred))
plt.title("Test data - Confusion Matrix")
sns.heatmap(cm_te, ax=ax_te, fmt='d', cmap='YlGnBu', annot=True)
# labels, title and ticks
ax_te.set_xlabel('Predicted labels');
ax_te.set_ylabel('True labels');
ax_te.set_ylim(2.0, 0)
ax_te.xaxis.set_ticklabels(['No', 'Yes']);
ax_te.yaxis.set_ticklabels(['No', 'Yes']);

plt.show()
return

```

```

[44]: def draw_roccurve(y_tr,y_tr_pred,y_te,y_te_pred):
    #fpr,tpr,thresholds
    fpr, tpr, thresholds = roc_curve(y_tr, np.array(y_tr_pred))
    #auc score train score
    auc_train = round(auc(fpr, tpr),5)
    plt.plot(fpr, tpr, label=" AUC train "+str(auc_train))
    plt.plot([0, 1], [0, 1], 'r--')

    fpr, tpr, thresholds = roc_curve(y_te, np.array(y_te_pred))
    #auc score test score
    auc_test = round(auc(fpr, tpr),5)
    plt.plot(fpr, tpr, label=" AUC test "+str(auc_test))

```

```

plt.plot([0, 1], [0, 1], 'b--')

plt.legend()
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.title("ROC" )
plt.grid()
plt.show()
return auc_train, auc_test

```

```

[52]: from sklearn.utils import class_weight
#class_weights = class_weight.compute_class_weight('balanced', np.
    ↳unique(y_train), y_train)
#class_weights = dict(zip(np.unique(y_train), class_weight.
    ↳compute_class_weight('balanced', np.unique(y_train), y_train)))
#y_integers = np.argmax(y_train, axis=1)
class_weights = class_weight.compute_class_weight(class_weight='balanced',
    ↳classes=[0,1], y=y_train)
d_class_weights = dict(enumerate(class_weights))
d_class_weights

```

```

[52]: {0: 0.5439099467262235, 1: 6.193470811038297}

```

```

[53]: #based on the best parameters, predict values and plot AUC and return the model
#def measure_accuracy(study,X_tr,X_te,y_tr,y_te):
def measure_accuracy(X_tr,X_te,y_tr,y_te):

    #clf = XGBClassifier(**study.best_params) # overfitting with optuna results.
    ↳ Hence manually changed the params as below.
    param_grid = {
        "verbosity": 0,
        "objective" : 'binary:logistic',
        "n_estimators":10000,
        "learning_rate": 0.001,
        "alpha": 0.001,
        "gamma": 0.005,
        "lambda": 0.004,
        "min_child_weight": 20,
        'subsample' : 0.8,
        'colsample_bytree': 0.9,
        'colsample_bylevel': 0.5,
        "max_depth": 7,
    }

    clf = XGBClassifier(tree_method = "exact", predictor =
    ↳ "cpu_predictor", sample_weight=d_class_weights, random_state=0)
    clf.set_params(**param_grid)

```

```

clf.fit(X_tr,y_tr,verbose=False, )

y_tr_pred = batch_predict(clf, X_tr)
y_te_pred = batch_predict(clf, X_te)

plot_confusionmatrix(y_tr,y_tr_pred,y_te,y_te_pred)
print('='*70)
auc_train,auc_test=draw_roccurve(y_tr,y_tr_pred,y_te,y_te_pred)
print('='*70)
return clf, auc_train,auc_test


```

```

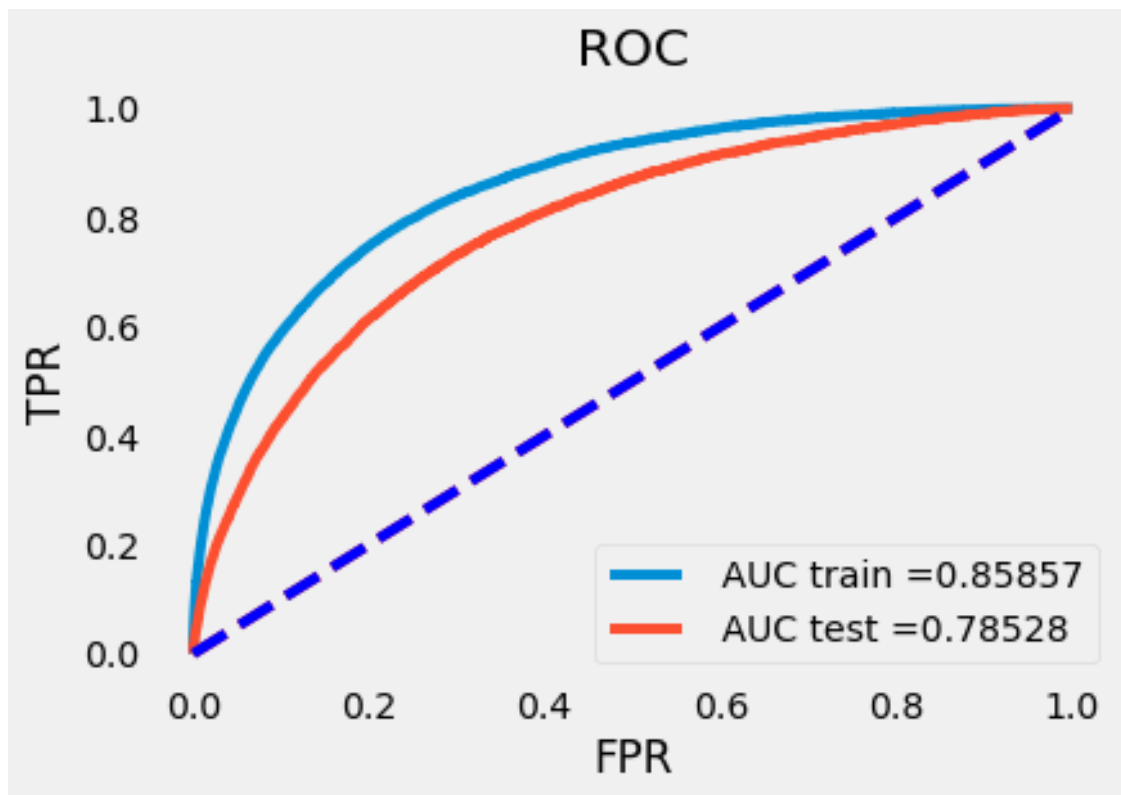
[54]: #xgb_model,auc_tr_xgb_model,auc_te_xgb_model =
      ↪measure_accuracy(study,X_train,X_test,y_train,y_test)
xgb_model,auc_tr_xgb_model,auc_te_xgb_model =
      ↪measure_accuracy(X_train,X_test,y_train,y_test)

```

Training data tn-> 189201, fp-> 198, fn-> 15525, tp-> 1108  
 Training data tn-> 93093, fp-> 194, fn-> 7915, tp-> 277



=====




```

=====
[55]: top20_feature_names=[]
      feature_importance = xgb_model.feature_importances_
      feature_importances = (xgb_model.feature_importances_ / sum(xgb_model.
      ↪feature_importances_+0.000001)) * 100
      indices = feature_importance.argsort()[::-1][:20]
      for i in indices:
          top20_feature_names.append(feature_names[i])

      #Plot bar plot for top 15 features
      plt.close()
      column =top20_feature_names
      score = feature_importance[indices]
      plt.figure(figsize =(25, 8))
      sns.barplot(x=column, y=score, palette="viridis")
      plt.xticks(rotation=90)
      plt.title('Top 20 Features')
      plt.show()

```





```
[56]: # http://zetcode.com/python/prettytable/
from prettytable import PrettyTable

x = PrettyTable()
x.field_names = ["Feature", "Score"]
for val in zip(column,score):
    x.add_row([val[0],np.round(val[1],4)])

x.sortby = "Score"
x.reversesort = True
x.align["Feature"] = "l"
x.align["Score"] = "r"

print(x)
```

```
+-----+-----+
| Feature                                | Score |
+-----+-----+
| EXT_SOURCE_2 EXT_SOURCE_3             | 0.0144 |
| EXT_SOURCE_1 EXT_SOURCE_3             | 0.0064 |
| EXT_SOURCE_1 EXT_SOURCE_2             | 0.0064 |
| EXT_SOURCE_2 Age                       | 0.0058 |
| EXT_SOURCE_3^2                         | 0.0052 |
| EXT_SOURCE_3_x                         | 0.005  |
| EXT_SOURCE_3_y                         | 0.0047 |
| EXT_SOURCE_2_y                         | 0.0034 |
| EXT_SOURCE_2_x                         | 0.0033 |
| NAME_EDUCATION_TYPE_1                  | 0.0031 |
| NAME_EDUCATION_TYPE_0                  | 0.0031 |
```

|                                       |        |
|---------------------------------------|--------|
| EXT_SOURCE_2^2                        | 0.0031 |
| bureau_AMT_CREDIT_SUM_OVERDUE_mean    | 0.0028 |
| bureau_AMT_CREDIT_SUM_OVERDUE_max     | 0.0028 |
| bureau_CREDIT_TYPE_Microloan_count    | 0.0027 |
| EXT_SOURCE_3 Age                      | 0.0027 |
| pre_app_RATE_INTEREST_PRIVILEGED_mean | 0.0026 |
| bureau_AMT_CREDIT_SUM_OVERDUE_sum     | 0.0026 |
| OWN_CAR_AGE                           | 0.0026 |
| CODE_GENDER_1                         | 0.0026 |

+-----+-----+

```
[57]: import joblib
      # save model
      joblib.dump(xgb_model, './results/model_xgboost.pkl')
```

```
[57]: ['./results/model_xgboost.pkl']
```